

Supervised by: Dr. Farouk Mansouri, LuxProvide. Prof. Olaf Schenk, USI. **Project:** Benchmarking AI Factories on MeluXina supercomputer. **Challenge:** Develop a comprehensive benchmarking framework for AI infrastructure evaluation including LLM inference, vector databases, and storage systems. **Course:** Software Atelier: Simulation, Data Science & Supercomputing SA 2025-2026.

Keywords: High-Performance Computing, AI Benchmarking, LLM Inference, Vector Databases, Prometheus, Grafana, SLURM, Apptainer, GPU Acceleration.

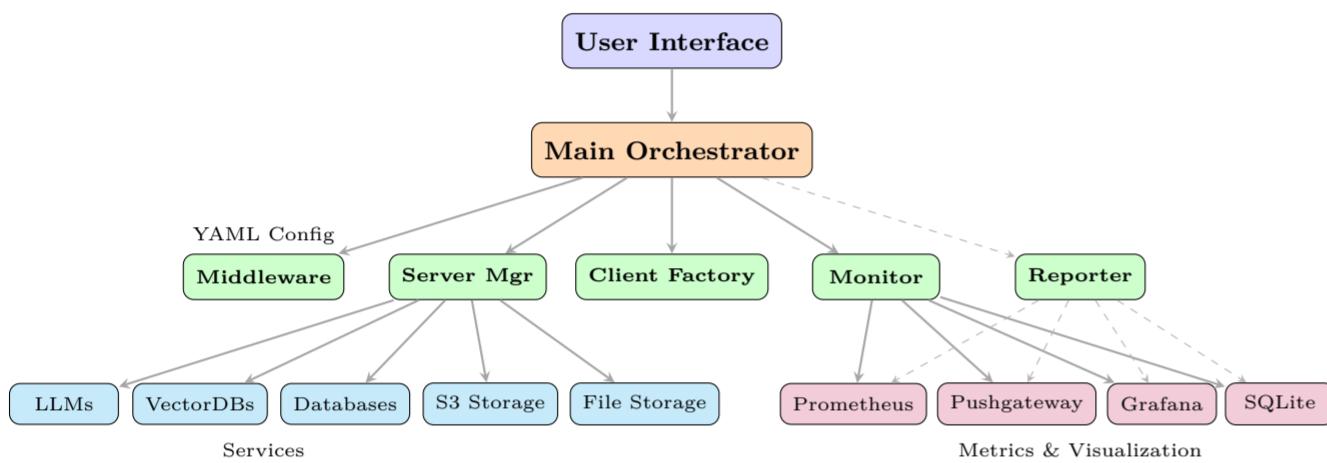
Motivation and Objectives

The emergence of **AI Factories** requires robust evaluation methodologies for AI infrastructure components. This project addresses the critical need for:

- ❶ **Unified benchmarking** across diverse AI services.
- ❷ **Real-time monitoring** with production-grade tools.
- ❸ **HPC integration** for large-scale evaluation.
- ❹ **Reproducible results** via modular architecture.

System Architecture

The framework consists of five core components orchestrated by a main controller:



Supported Services

The framework supports comprehensive benchmarking of AI Factory components:

Service Category	Backends	GPU
LLM Inference	Ollama, vLLM, Triton	✓
Vector Databases	ChromaDB, Faiss, Weaviate	✓
Relational DB	PostgreSQL	✓
Object Storage	MinIO (S3-compatible)	✓
File Storage	POSIX Filesystem	✓

All services support GPU acceleration via NVIDIA containers with `--nv` flag.

Benchmark Configuration

YAML-based recipe configuration enables reproducible benchmarks:

```

recipe.yml
benchmark:
  name: "AI Factory Benchmark"
  duration: 600
  services:
    - service.name: ollama-llama2
      service.type: ollama
      container.image: docker://ollama/ollama
      port: 11434
      client.count: 5
      requests.per.second: 10
      slurm:
        partition: gpu
        time: "01:00:00"
  
```

Execution Modes:

- **Parallel:** Concurrent multi-service requests through ThreadPoolExecutor.
- **Sequential:** For debugging and resource-constrained runs.
- **Interactive:** Via SLURM `salloc` sessions.

Implementation Details

Technology Stack

- ❶ **Orchestration:** Python 3.12, SLURM.
- ❷ **Containerization:** Apptainer.
- ❸ **Monitoring:** Prometheus, Grafana.
- ❹ **Storage:** SQLite, JSON reports.
- ❺ **GPU:** NVIDIA with `--nv`.



Dependencies: PyYAML, requests, numpy, pandas, psycopg2-binary, boto3, botocore, chromadb, faiss-cpu, pymilvus, weaviate-client, prometheus-client, rich, click, tqdm, matplotlib, seaborn, plotly.

Workflow and Execution

Execution Flow:

1. Load recipe (YAML).
2. Initialize logger/monitor.
3. Start SLURM services.
4. Execute benchmarks.
5. Collect metrics.
6. Generate reports.
7. Push to Pushgateway.
8. Cleanup services.

SLURM: Auto batch scripts, dynamic endpoints, GPU support.

Outlook

- Multi-node distributed benchmarks.
- Automated scaling analysis studies.
- ML training workload benchmarks.
- Network bandwidth/latency analysis.

Key influential references

- [1] LuxProvide MeluXina Supercomputer. [2] Prometheus Monitoring. [3] Grafana Dashboards. [4] SLURM Workload Manager. [5] Apptainer Containers. [6] Ollama LLM [7] vLLM Inference. [8] NVIDIA Triton. [9] ChromaDB Vector DB. [10] Faiss Vector Search. [11] Weaviate Vector DB. [12] PostgreSQL Database. [13] MinIO Object Storage.